

ENEE 467: Robotics Project Laboratory

Hardware-Based 3D Perception, Motion Planning, and Control for the UR3e Robot

Lab 7

Learning Objectives

The purpose of this lab is to provide a hands-on introduction to 3D perception, motion planning, and control on physical hardware. Students will work with point-cloud-based pose estimation using ROS 2 and the Intel RealSense D435i RGB-D camera, and perform motion planning and control on the UR3e robot using MoveIt 2 and `ros2_control`. After completing this lab, students should be able to:

- Integrate point cloud-based visual perception with motion planning
- Plan and execute manipulation trajectories based on visual feedback
- Understand the perception-to-action pipeline in ROS 2

Autonomous robotic manipulation unites perception, planning, and control to enable robots sense, reason, and act independently towards the realization of a prescribed task. Among these elements, this lab focuses on perception and control to enable the real UR3e robot to detect and plan motions to an object in its workspace from point cloud input. Through this lab, students will explore how these subsystems interact in real time, in simulation and partly in hardware.

Lab Instructions

Follow the lab procedure closely. Some code blocks have been populated for you and some you will have to fill in yourself. Pay close attention to the lab [Questions](#) (highlighted with a box). Answer these as you proceed through the lab and include your answers in your lab writeup. There are several lab [Exercises](#) at the end of the lab manual. Complete the exercises during the lab and include your solutions in your lab writeup.

Autonomous Object Relocation

While conceptually simple, autonomous *object relocation*, or more commonly, pick-and-place (PnP), integrates all aspects of robotic manipulation, viz:

1. **Perception:** Detect and estimate the pose of the object in 3D space.
2. **Motion Planning:** Generate a collision-free trajectory to reach the object, execute the grasp, and move to the goal pose.
3. **Execution:** Command the manipulator to follow the planned motions and monitor execution for success or failure.

List of Questions

Question 1.	4
Question 2.	9
Question 3.	11
Question 4.	11

List of Exercises

Exercise 1. Evaluating Autonomous Object Relocation in Simulation (35 pts.)	4
Exercise 2. Simple Pick-and-Place Action Client (35 Pts.)	12

Hardware Safety

Throughout this manual, boxes beginning with a **Question #** label contain questions you must answer as you progress through the lab:

Question 0. A question for you to consider as you work.

⚠ Safety and Setup Notices. Warnings highlight critical steps for hardware safety, teach-pendant operation, and common setup pitfalls. Always read these carefully before commanding motion on the UR3e.

⚠ Keep clear of the robot workspace when running autonomous motion. Confirm that the pendant is in Play mode before proceeding.

📖 Notes. These boxes provide background or theoretical context for the UR robot system. They are optional during lab time but useful for your report:

📖 UR3e kinematics are defined in a six-joint chain; MoveIt uses the URDF and joint limits to compute feasible IK solutions.

Code, commands, and output are formatted as follows:

```
ros2 launch ur_robot_driver ur_control.launch.py ...
```

```
[INFO] [launch]: UR3e driver successfully connected.
```

Files and directories appear in magenta teletype font (e.g., `~/lab7_ws/src`), while in-text code and variables use blue teletype (e.g., `object_pose`). Python type hints or XML tags appear in green (e.g., `str`, `<param>`).

1. Pre-Lab: Run the Simulation Baseline (No Coding Required)

Before diving into the hardware, you will first observe a baseline simulation that demonstrates how perception, planning, and control integrate in an environment closely mirroring the real hardware. Then, you'll record task metrics for your final presentation. This unit requires no code modifications. Confirm that you have the simulator Docker image by running the following command:

```
docker image ls | grep lab-7-sim
```

Confirm that **lab-7-sim-image** is printed to the terminal. Then run the following command:

```
cd ~/Labs/lab-7-sim/docker
```

Next, start the **pre-installed** simulator Docker container using:

```
userid=$(id -u) groupid=$(id -g) docker compose -f lab-7-sim-compose.yml \
run --rm lab-7-sim-docker
```

Because this unit is simulation-only and does not require code editing, we do not mount any host folders into the container. All required packages are pre-installed inside the image, and you will work entirely within the container environment.

1. In the Docker container, source your ROS 2 workspace:

```
cd ~/ros2_ws
colcon build --symlink-install
source install/setup.bash
```

2. Launch the simulated UR3e-Hand-E robot with the perception, planning, and control nodes:

```
ros2 launch ur3e_hande_bringup sim.launch.py
```

3. Verify that separate Gazebo and RViz windows appear displaying the following:

- The UR3e-Hand-E robot model loaded in Gazebo.
- Segmented point cloud and plane RViz markers for the table and object (red cylinder). Note that the red cylinder is a proxy for the “cup” in the real environment. Simulating grasping in Gazebo has been a longstanding [issue](#).
- The point cloud based object pose action server should also display a prompt indicating successful detection and publishing of several objects

4. Then in a new `tmux` pane, start the pick-and-place node:

```
ros2 run ur3e_hande_moveit_py pnp_demo_sim
```

If everything functions correctly, you should now observe a complete pick-and-place cycle executed autonomously in simulation. The pick-and-place node will also print output to the terminal reporting successful task execution for each object relocation sub-task (the following will only be recorded at the END of the task). See [Figure 1](#) for a screenshot of the robot at various stages of the task.

```
Execution completed: SUCCEEDED
Retreated to home; closing gripper.
```

```
Closing gripper...
Pick-and-place sequence complete.
```

Question 1. Which ROS 2 nodes launch during the simulation test above? Enumerate these nodes, and classify the nodes in your list into **Perception**, **Planning**, and **Control** using a table. (**Hint:** Run: `ros2 node list`, while the simulation is active. The node names are descriptive enough to identify their respective classes. Exclude RViz nodes.)

Exercise 1. Evaluating Autonomous Object Relocation in Simulation (35 pts.)

In this exercise, you will sample an autonomous object relocation node and collect data under varying joint velocities. **Note:** We initially planned that data from this part would serve as a benchmark for the hardware variant of the object relocation task, but due to circumstances beyond our control, the Hand-E grippers are not able to be installed at this time, and therefore will not be available for use in the hardware procedure.

Deliverables:

Complete this exercise by performing the following steps:

Video/Frame Screenshots:

- a. Include a **single** video OR a set of screenshots showing:
 - Object detection and plane segmentation results from your RViz window
 - The robot at the approach and place configurations in Gazebo (see [Figure 1b](#) and [d](#), respectively).

Include the video clip in your Canvas submission with your report. Alternatively, you can include a collage of screenshots at each task phase in your report.

Planning Metrics

- b. Terminate all running ROS 2 processes, then start the simulator setting the `pnP` and `print_metrics` parameters equal to `true`:

```
ros2 launch ur3e_hande_bringup sim.launch.py pnp:=true \
print_metrics:=true max_vel_scale:=0.25
```

where `max_vel_scale` is a factor that scales the velocity of each joint of the arm. Once the task is completed, the following metrics will be printed to shell (**Note:** your numbers may be slightly different, but the entire PnP task should succeed):

```
-----
Metrics Summary:
1. Planning Success (bool): 1
2. Planning time [s]: 0.0085 s
3. Execution time [s]: 4.0520 s
4. Time (PnP Pipeline) [s]: 34.3 s
-----
```

Record the printed metrics in your report.

- c. Execute the simulation command for **three** new runs setting `max_vel_scale` to 0.20, 0.15, and 0.1, in each case. Copy the metrics directly into your report for later number-crunching and analysis (d. and e.)

[DO THIS AFTER THE LAB SESSION]

Post-Lab Analysis and Discussion

- d. Summarize the metrics you recorded for all runs (three runs in (c) and the run in (b) with `max_vel_scale` = 0.25, making four altogether) in a well-formatted table. Briefly discuss any observed trends. For example:
- How does reducing velocity scaling affect execution time?
 - Does planning time change significantly?
 - Does success remain consistent across runs with changing max joint velocity?
- e. Compute the average planning and execution times, and provide these values in your report.




Figure 1: Phases of the object relocation task visualized in simulation.

After completing the steps above (except d. and e.), proceed with the rest of the lab procedure.

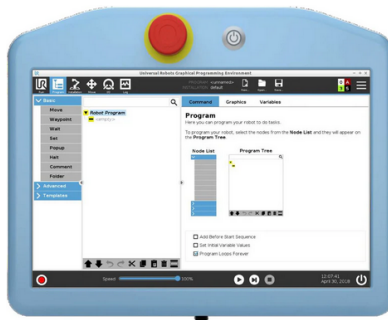
THE NEXT UNIT USES A DIFFERENT DOCKER IMAGE. ENSURE YOU RUN THE COMMANDS HEREFTER WITHIN A CONTAINER CREATED FROM THE HARDWARE IMAGE (`lab-7-hw-image`).

2. Hardware Lab Procedure

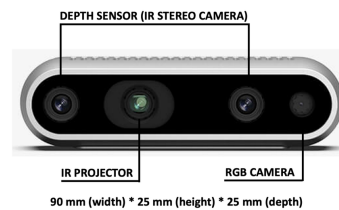
⚠ Stay alert when working with the robot, and use the Emergency Stop Button on the Teach Pendant (see [Figure 2a](#)) if needed. Always inform the team member(s) at the hardware station before sending commands to the robot.

2.1 Prerequisites

This lab assumes familiarity with ROS 2, MoveIt, and Universal Robotics's (UR) robot drivers and hardware interface. Prior experience with the teach pendant is beneficial but not required.



(a) UR3e Teach Pendant. Emergency Stop button in red.



(b) RealSense D435i Stereo RGB-D camera.

Figure 2: Teach Pendant (TP) and RealSense D435i camera.

2.2 Getting the Lab Code

From your host system, open a terminal, **cd** to the **~/Labs** folder, and clone the lab repository:

```
cd ~/Labs
git clone https://github.com/ENEE467-F2025/lab-7.git
cd lab-7/docker
```

Verify that the Docker image has built successfully:

```
docker image ls
```

You should see output similar to:

```
lab-7-hw-image    latest    <container_id>    <time_since_inst>    <size>B
```

Launch the Docker container:

```
docker compose -f lab-7-hw-compose.yml run --rm lab-7-hw-docker
```

You should now see:

```
(lab-7) robot@docker-desktop:~$
```

2.3 Building the Workspace

Inside the container, build and source your ROS 2 workspace:

```
cd ~/ros2_ws
colcon build --symlink-install
source install/setup.bash
```

2.4 Hardware Setup & Verification

Before starting the hardware procedure, verify that all components of the UR3e system are functional by performing the following steps:

A. RealSense D435i Stereo RGB-D Camera:

- From within a running container, run the following command:

```
realsense-viewer
```

You should see the following window. Confirm that the USB standard (shown on the top-left corner of the app) is at least 3.0 (the lab computers have 3.2).

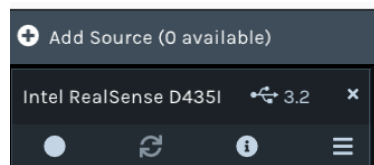


Figure 3: Partial view of the Intel RealSense viewer highlighting the USB standard.

- Close the RealSense Viewer program by pressing **Ctrl + C** in the previous tmux pane. Then run the following command to verify that your camera is configured correctly with ROS 2:

```
ros2 launch realsense2_camera rs_pointcloud_launch.py
```

If you see an RViz window with color, depth, and point cloud streams, then your camera has been setup correctly, and you can terminate the camera node and proceed with the rest of the procedure.

⚠ If you see a lower USB standard, close the viewer, **carefully** unplug the RealSense camera (at the camera end and not the PC end), plug it back, and try running the viewer again.

B. UR3e Robot Arm:

- Power on the TP using the button beside the red emergency stop. Once the PolyScope GUI loads, the robot will be in a powered-off state. To enable it, click the bottom-left red button (next to **Power off**), then press ON to switch the robot to the **IDLE** state.
- Verify that the robot is reachable over Ethernet by running the following command from your container session:

```
ping 192.168.77.22
```

You should see output printed to the terminal indicating successful communication:

```
64 bytes from 192.168.77.22: icmp_seq=49 ttl=64 time=0.329 ms
```

C. Workspace Props:

- Verify that the object (cup) we will use in the pose estimation portion of the procedure is available in the robot's workspace (**Note:** cup colors will vary, but this information is immaterial to the task)

2.5 Launching MoveIt 2 for the Arm-Gripper System

Before controlling the robot via ROS 2, we need to start the TP UR3e program. Power ON the TP, if it isn't already powered on, then:

1. Click Power Off → ON at the bottom-left
2. Step out of the robot's workspace, then click the → START button to unlock the robot's brakes and allow it receive programs. You should now see to [Manual](#) on the bottom left.
3. Tap [Exit](#) to return to the main GUI. Then tap Load Program, and select `ur3e_ros.urp`.
4. Assuming packages in your workspace were built with `--symlink-install`, in a sourced `tmux` session, split panes. In the first pane, start the UR3e ROS 2 hardware driver (`ur_robot_driver`):

```
ros2 launch ur_robot_driver ur_control.launch.py \  
ur_type:=ur3e robot_ip:=192.168.77.22 launch_rviz:=false
```

5. Click Open then press [Play](#) to start the program.

The lab computers have been setup to automatically read the configuration file passed as the `kinematics_params_file`, so you do not need to provide those yourself.

⚠ NOTE: The pendant must remain in [Play](#) mode for ROS 2 to control the robot. Only click the [Play](#) button when you are ready to command the robot via ROS 2. The robot must remain in the [IDLE](#) or [OFF](#) states at all other times.

2.5.1 Plan Motions on the UR3e using MoveIt

MoveIt 2 provides IK, collision checking, trajectory generation, and controller management for the UR3e. Turn ON the TP and start the UR3e ROS 2 driver, if you haven't already. Click [Play](#) on the TP to ready the robot for programming.

Then, in a new `tmux` pane, launch the MoveIt bringup for the UR3e:

```
ros2 launch ur_moveit_config ur_moveit.launch.py \  
ur_type:=ur3e robot_ip:=192.168.77.22
```

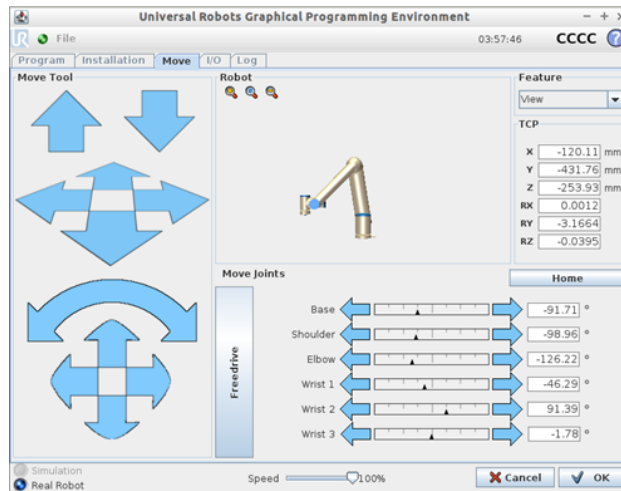



Figure 4: Robot jogging interface on the TP.

RViz will open and load the UR3e's URDF. Using the Move tab on the TP, jog the real robot SLOWLY, and confirm RViz mirrors the motion. If joint states do not update, MoveIt cannot plan safely. Then, maintaining a safe distance from the UR3e, using MoveIt's [MotionPlanning](#) RViz panel, plan then execute a motion to the `ur_test` configuration by selecting it from the Goal State dropdown. You should observe the UR3e move to the selected configuration, with its state evolution mirrored on RViz.

Question 2.

The UR3e ROS 2 driver publishes joint states on the topic: `/joint_states`.

- Why must MoveIt receive these messages continuously before planning?
- Terminate the UR3e ROS 2 driver process and try planning to a goal from MoveIt. What failure mode, if any, do you record?

2.5.2 Programmatic Control via PyMoveIT2

Similar to Lab 5, Part 2, we can send programmatic goals to the robot by scripting. Open the `ur3_joint_goal.py` file in the `ur3e_hande_moveit_py` ament_python package. Ensure that both the UR3e drivers and MoveIt are active as per the commands introduced earlier. After notifying the team member(s) manning the hardware station to keep a safe distance, the team member(s) sending commands can run:

```
ros2 run ur3e_hande_moveit_py ur3_joint_goal
```

You should observe the robot move to the test configuration earlier. You can home the robot by specifying its named group state within the `group_state` parameter as follows:

```
ros2 run ur3e_hande_moveit_py ur3_joint_goal --ros-args -p \
group_state:="ur_home"
```

3. Point Cloud-Based Object Pose Estimation

3.1 Test the Perception Nodes

The RealSense D435i driver (`realsense2_camera`) publishes several visual message streams:

- `/camera/camera/color/image_raw`: RGB image
- `/camera/camera/depth/image_rect_raw`: registered depth
- `/camera/camera/color/camera_info`: intrinsic calibration
- `/camera/camera/depth/color/points`: aligned organized point cloud

Terminate all running ROS 2 processes. Then bring up the camera and point cloud perception nodes using the following command (ensure the cup is in the camera's field-of-view):

```
ros2 launch ur3e_hande_perception perception.launch.py
```

The pipeline detects an object by clustering point clouds and fitting clusters to a prespecified known geometry within bounds. The command above starts nodes that perform:

- **Voxel filtering**: reducing point cloud density
- **Cropping**: removing irrelevant workspace regions
- **Plane segmentation**: extracting tabletop surfaces
- **Object clustering**: grouping and bounding-boxing objects
- **6D pose extraction**: estimating object centroids and rough extents

An RViz window will pop-up displaying an RGB image overlay, a filtered cloud, plane markers (green), and object bounding box (cyan colored here, but object colors might vary; see [Figure 5](#)).

🍃 If the perception node detects multiple objects in the workspace or no object at all, you will need to adjust the `crop_bounds` parameter used by the `pc_voxel_filter_node`. This parameter can be found in the `perception.launch.py` file located in the `launch` directory of the `ur3e_hande_perception` ament_python package. Modify the `x` and `y` bounds (either increasing or decreasing them based on your situation), then re-run `perception.launch.py`.

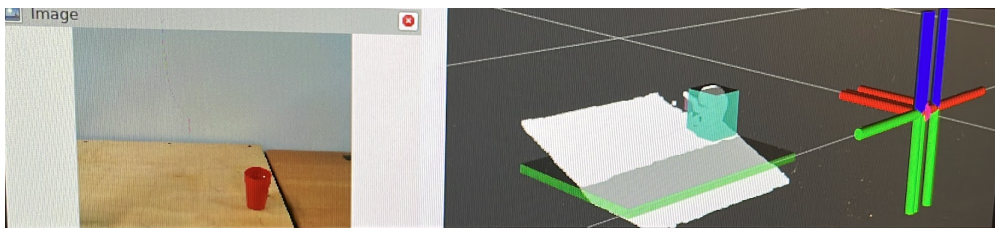


Figure 5: RViz window at camera-only bringup showing camera frames and RViz box markers highlighting the detected object (cup) and segmented plane (table).

We can check the pose of the detected object that provides the input for grasping and place pose computation by listening to the following topic:

```
ros2 topic echo --once /object_detected
```

If an object matching the provided dimensions is found within the camera's view, you should see the action server return its parameters should see its parameters:

```
object_id: 0
position:
x: -
y: -
z: -
dimensions: height: <> width: <> thickness: <>
```

The objects are detected by forming clusters from a filtered point cloud represented by the `PointCloud2` ROS 2 message type (see [PointCloud2 Message](#)) on the topic `/filtered_cloud`.

Question 3.

Try echoing a single stream from the `/filtered_cloud` topic. (Hint: use:

```
ros2 topic echo --once /filtered_cloud
```

- (a) What camera stream does the value of the `frame_id` in the `PointCloud2` message above relate to?
- (b) How does this frame relate to pose estimation? (**Hint**): see [From Meters to Pixels](#).

3.2 Querying for Poses via a Custom Action

While manually querying for poses via the CLI is snappy, it will not afford other ROS 2 nodes access to this information. A simple way of providing such access is via an action server. Start the perception nodes if they aren't running:

```
ros2 launch ur3e_hande_perception perception.launch.py
```

Then, in another tmux pane, manually request an *eye-balled* object position (replace the target position values below with the values from the yellow box above as well as values outside the robot's workspace, like $x \geq 1.25$, $y \geq 1.0$, and $z \geq 1.5$, and observe the result returned by the action server:

```
ros2 action send_goal /get_target_obj_pose \
ur3e_hande_planning_interfaces/action/GetTargetObjPose \
"{target_position: {x: 0.28, y: 0.50, z: 0.05}, \
target_height: 0.12, target_obj_bounds: [0.20, 0.33, 0.0, 0.0]}"
```

The parameter `target_obj_bounds` contains the approximate location of the XY-coordinate of the object on the UR3e's mounting table. You should see the object's position printed to the terminal in the former case as well as the action server reporting success, confirming that an object satisfying the above description exists within the camera's field of view.

Question 4.

- a. Which ROS 2 message types define the Goal, Result, and Feedback of the above action?
Hint: Run

```
ros2 interface show ur3e_hande_planning_interfaces/action/GetTargetObjPose
```

- b. Why is an action the preferred choice in this setting over say a service? In particular, what could go wrong in a PnP task if our robot's control node relied on a ROS 2 service for object poses?

Exercise 2. Simple Pick-and-Place Action Client (35 Pts.)

In this exercise, you will complete and extend the autonomous pick-and-place routine implemented in `pnp_demo.py`, located within the `ur3e_hande_moveit_pyament_python` package. This node serves as a ROS 2 action client that integrates motion planning, inverse kinematics, and gripper control into a full pick-and-place behavior. Currently, this code assumes that IK always succeeds, the gripper is always available, and placement positions are always valid. Obviously this is not always the case and your task will be to fix these assumptions and extend the PnP behavior.

Deliverables:

IK Computation for Each Sub-Task:

- a. Modify the `TODO` highlighted block to get the IK result for each sub-task using the `solve_ik` function:

```
self.approach_q = self.solve_ik(self.approach_pos, seed=home_q)
```

The task progression is as follows (gripper actions in blue):

Home → Approach → Pre-Grip → **Grasp** → Lift → Place → **Release** → Home

IK Failure Handling:

- b. The method `solve_ik()` may return `None`. Modify the code so that if any of `approach_q`, `pre_grip_q`, or `place_q` are `None`, the node logs an error and aborts the sequence safely.

Synchronous Execution:

- c. Setting the launch parameter `synchronous` equal to `True` ensures that the motion for each sub-task completes before the next begins (i.e., waits on `wait_until_executed()` after every `timed_motion()` call). What happens when `synchronous` is set to `False`?

Files to Submit:

Submit the following:

- A code snippet showing your added IK failure handling.
- A code snippet showing your IK logic

Note: Since the Hand-E grippers are currently unavailable, you will not be able to test your PnP routine on hardware. If time permits and the hardware becomes available, we may run your programs after the lab. Your grade for this exercise will be based solely on your submitted Python code and your answer to part (c).

4. Conclusion and Lab Report Instructions

After completing the lab, summarize the procedure and results into a well written lab report. Include your solutions to all of the question boxes in the lab report. Refer to the [List of Questions](#) to ensure you don't miss any. Include your full solution to the exercises in your lab report. If you wrote any Python programs during the exercises, include these files as a separate attachment, appending your team's number and last names (in alphabetical order) to the file, e.g., **Jane Doe, Alice Smith, and Richard Roe** in **Team 90** submitting their modified `main.py` would submit the file `main_Team90_Doe_Roe_Smith.py` and the report `Lab_<LabNumber>_Report_Team90_Doe_Roe_Smith.pdf`, where `<LabNumber>` is the number for that week's lab, e.g., **0, 1, 2**, etc. Submit your lab report along with any code to ELMS under the assignment for that week's lab. You must complete and submit your lab report by **Friday, 11:59 PM** of the week following the lab session. See the **Files** section on ELMS for a lab report template named `Lab_Report_Template.pdf`. The grading rubric is as follows:

Component	Points
Lab Report and Formatting	10
Question 1.	5
Question 2.	5
Question 3.	5
Question 4.	5
Exercise 1.	35
Exercise 2.	35
Total	100